



Hackathon part 2: Programming

<https://github.com/pmemhackathon/2019-01-23>

Andy Rudoff
pmem SW Architect, Intel

Essential Background

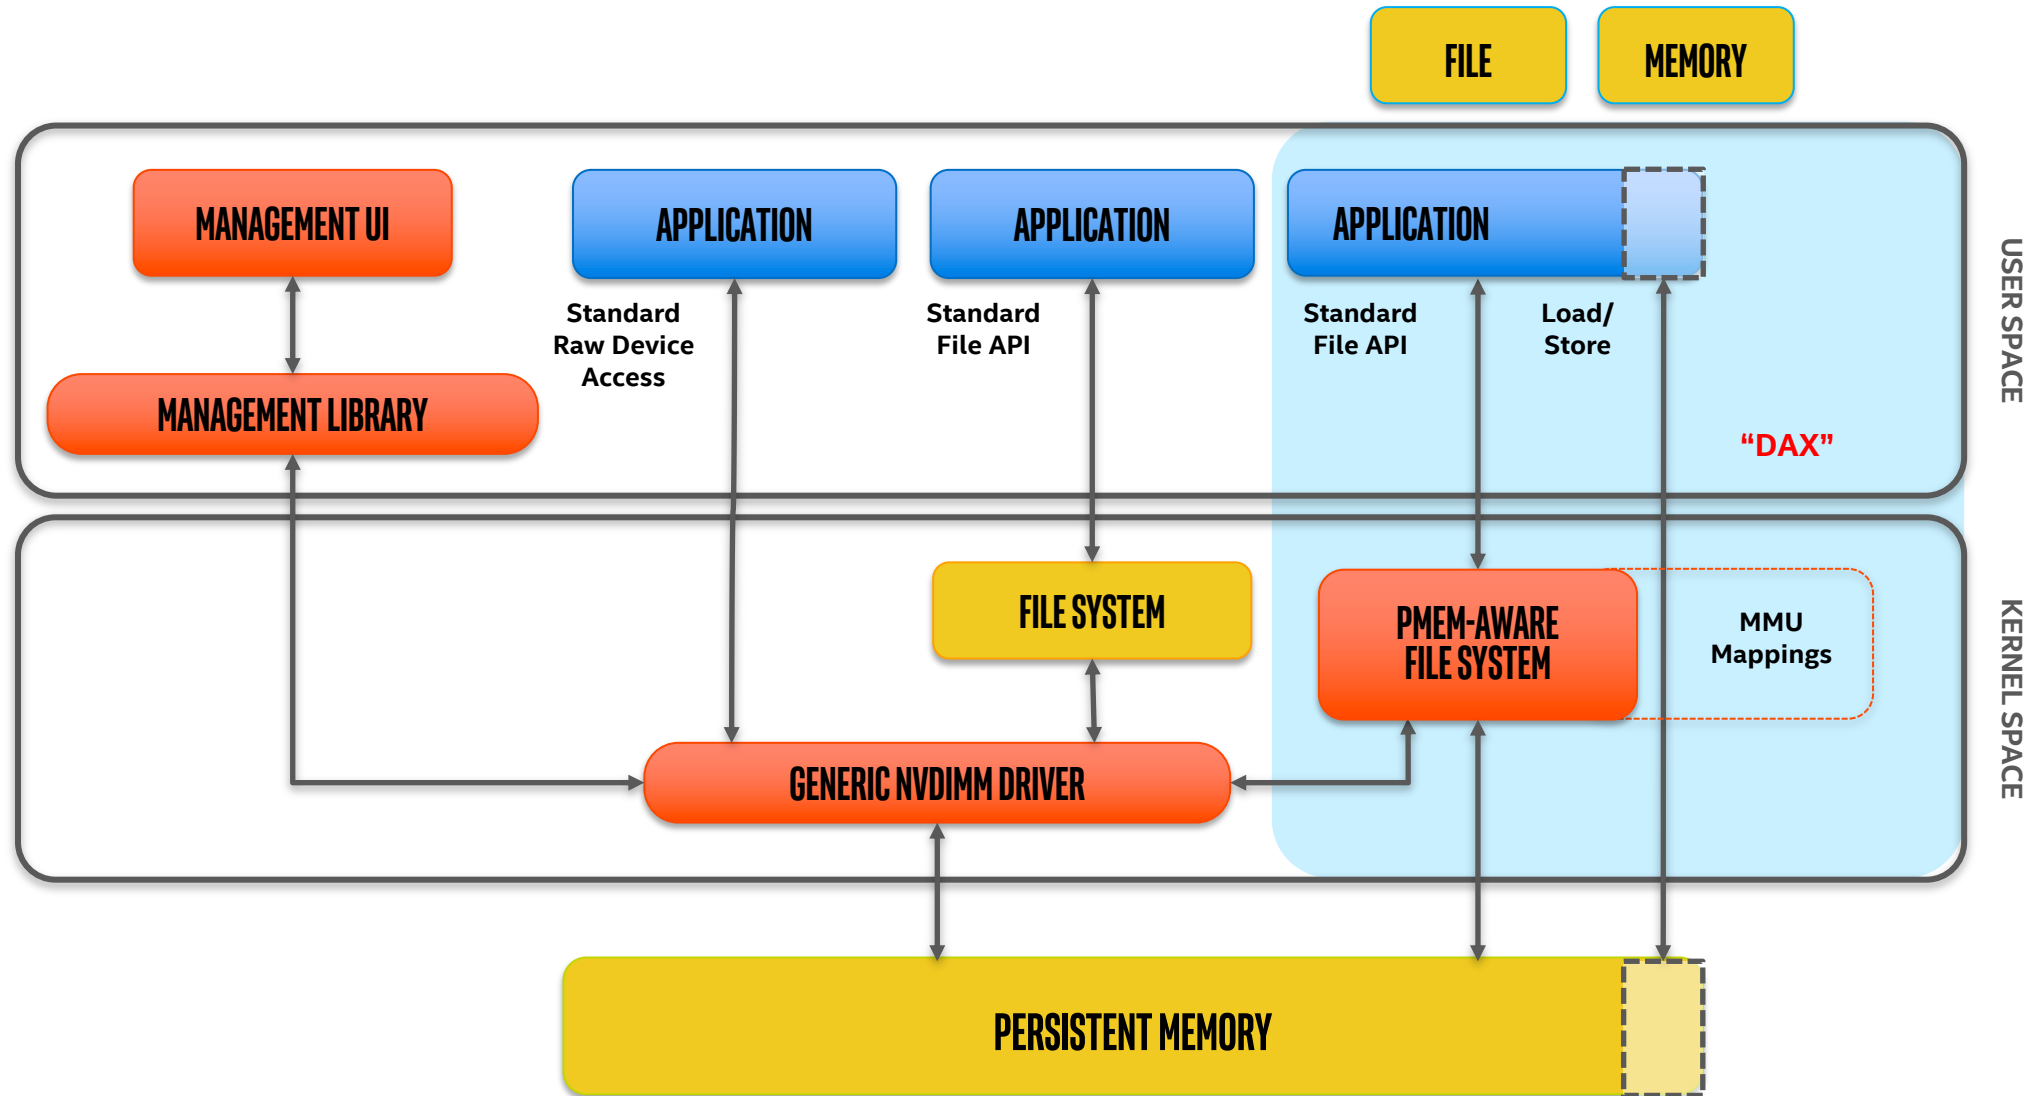
Lots of ways to use pmem with existing programs

- Storage APIs
- Libraries or kernels using pmem transparently
- Memory Mode

This hackathon doesn't cover the above (too easy!)

- We assume you want direct access to pmem
- We show code, but also concepts
- There are lots of paths you can take, these are just examples

The SNIA NVM Programming Model



Contents for this Hackathon

RAW Access to pmem

- `mmap()` -- you get a pointer to pmem, the rest is up to you
- Only 8-byte stores are powerfail atomic

libpmem

- One step above RAW access, still only 8-byte stores are powerfail atomic
- `mmap()`, `memcpy()` helper functions, optimized flush functions

libpmemblk

- Very simple transactional library, read/write fixed sized block only

libpmemobj

- **General-purpose allocations, transactions, atomics (series of examples)**

Pointers to related info:

- `libmemkind`, `libpmemkv`

Resources

- PMDK Resources:
 - Home: <https://pmem.io>
 - PMDK: <https://pmem.io/pmdk>
 - PMDK Source Code : <https://github.com/pmem/PMDK>
 - Google Group: <https://groups.google.com/forum/#!forum/pmem>
 - Intel Developer Zone: <https://software.intel.com/persistent-memory>
 - Memkind: <https://github.com/memkind/memkind> (see memkind_pmem(3))
 - libpmemkv: <https://github.com/pmem/pmemkv>
- NDCTL: <https://pmem.io/ndctl>
- SNIA NVM Programming Model: https://www.snia.org/tech_activities/standards/curr_standards/npm
- Getting Started Guides: <https://docs.pmem.io>

A Programmer's View (mapped files)

```
fd = open("/my/file", O_RDWR);  
...  
base = mmap(NULL, filesize,  
            PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0);  
close(fd);  
...  
base[100] = 'X';  
strcpy(base, "hello there");  
*structp = *base_structp;  
...
```

"Load/Store"

INSTALLING PACKAGES AND REPOS

Hackathon repo

<https://github.com/pmemhackathon/2019-01-23>

README.txt

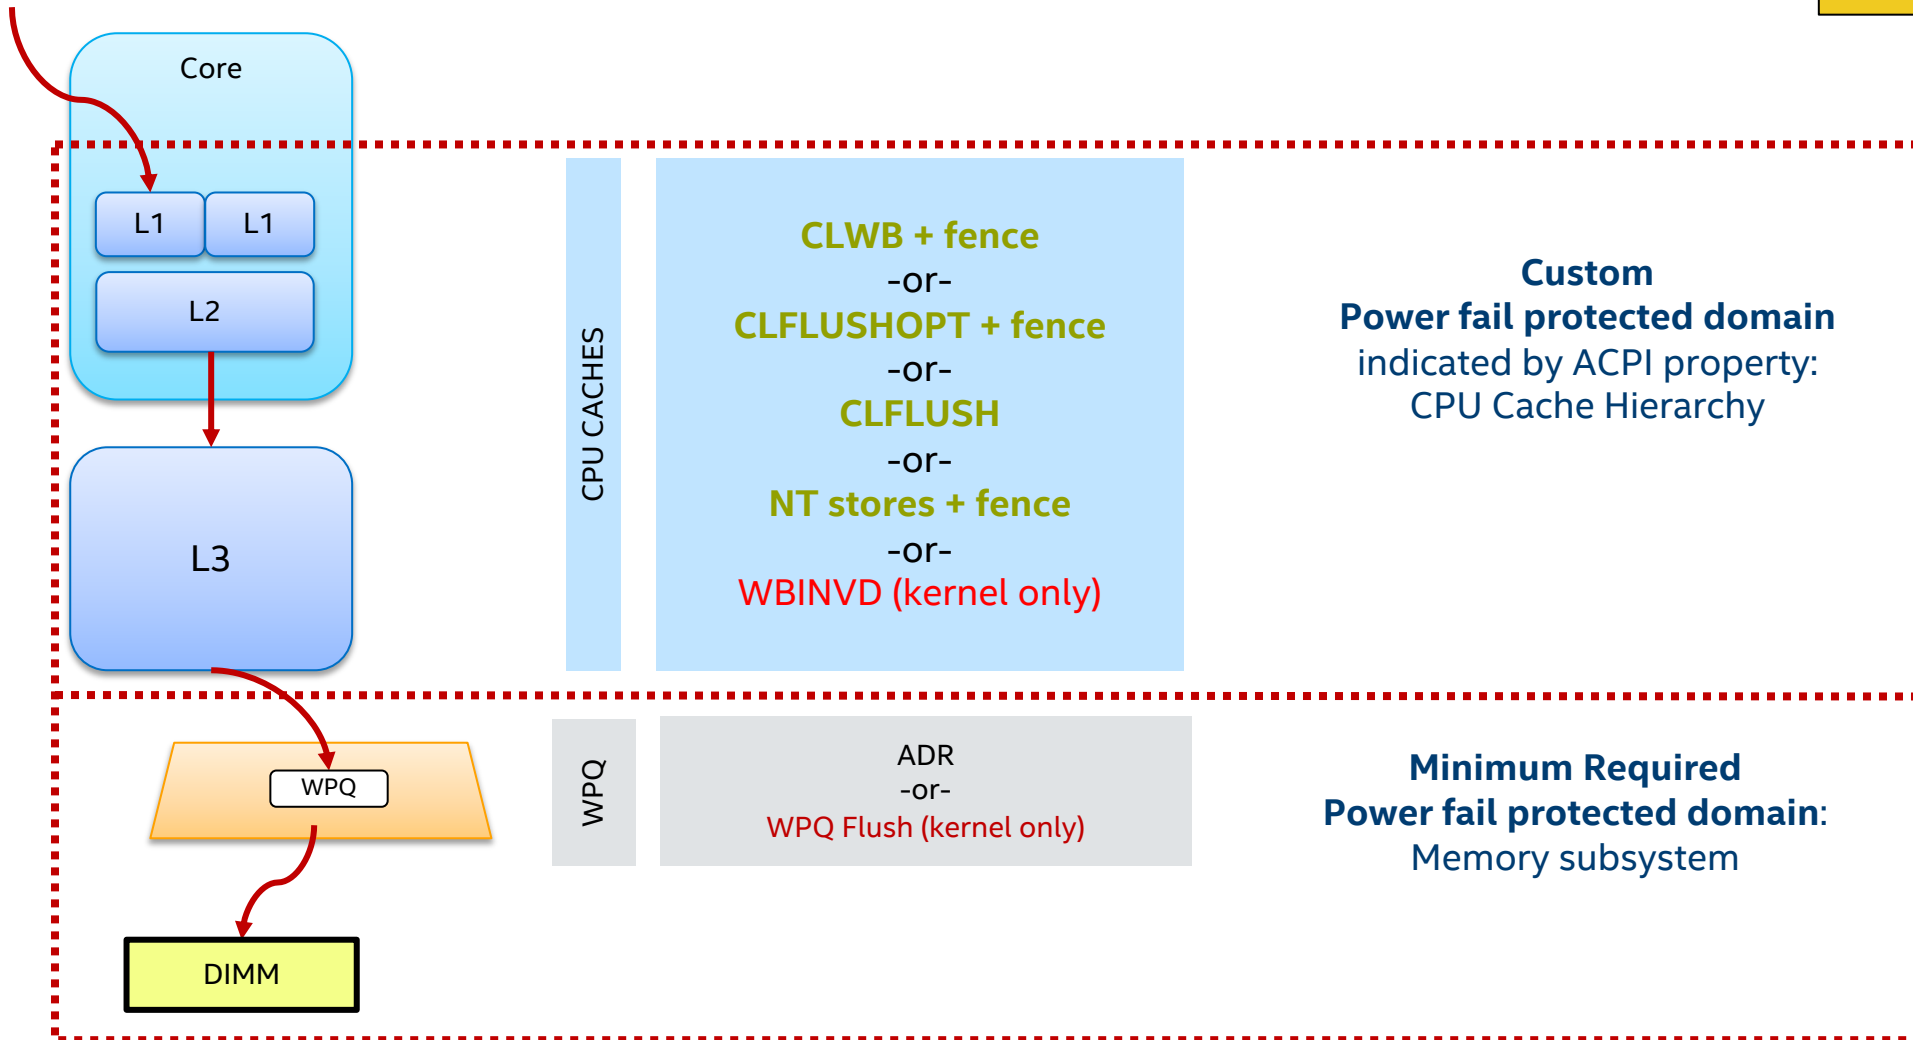
- We will walk through the examples using the README.txt
- Switch back to slides now and then for additional details
- Bringing up the README.txt in a window will help avoid typos
 - Cut and paste commands instead of reading and typing them

Repo contains these slides as well (slides.pdf) for future reference.

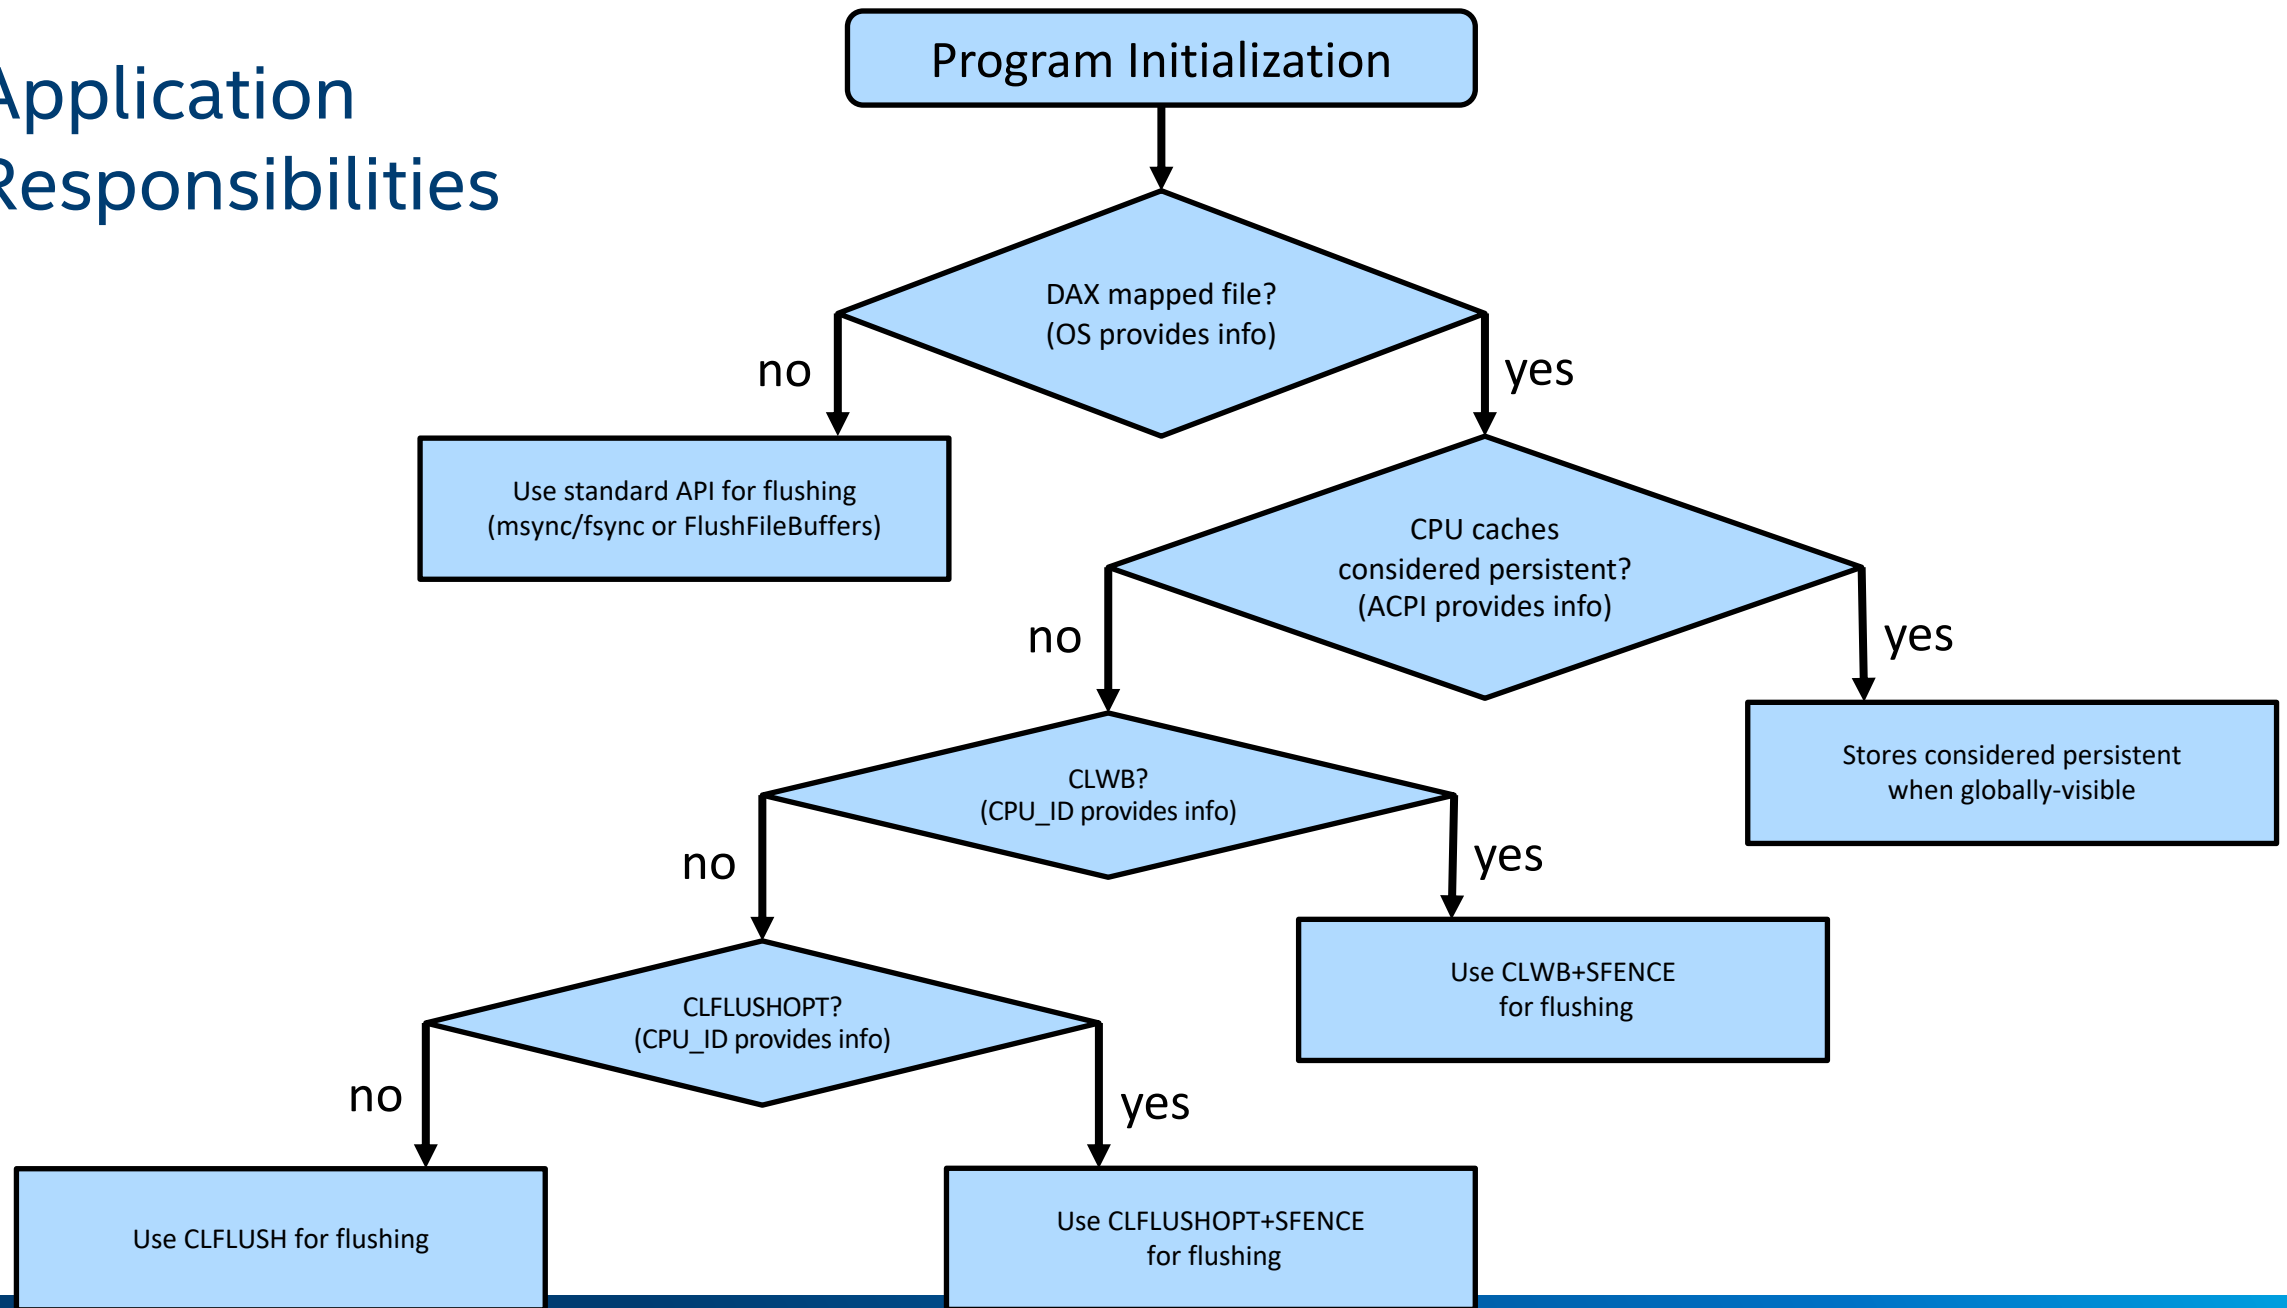
How the Hardware Works

MOV

Not shown:
MCA
ADR Failure Detection



Application Responsibilities



PMDK

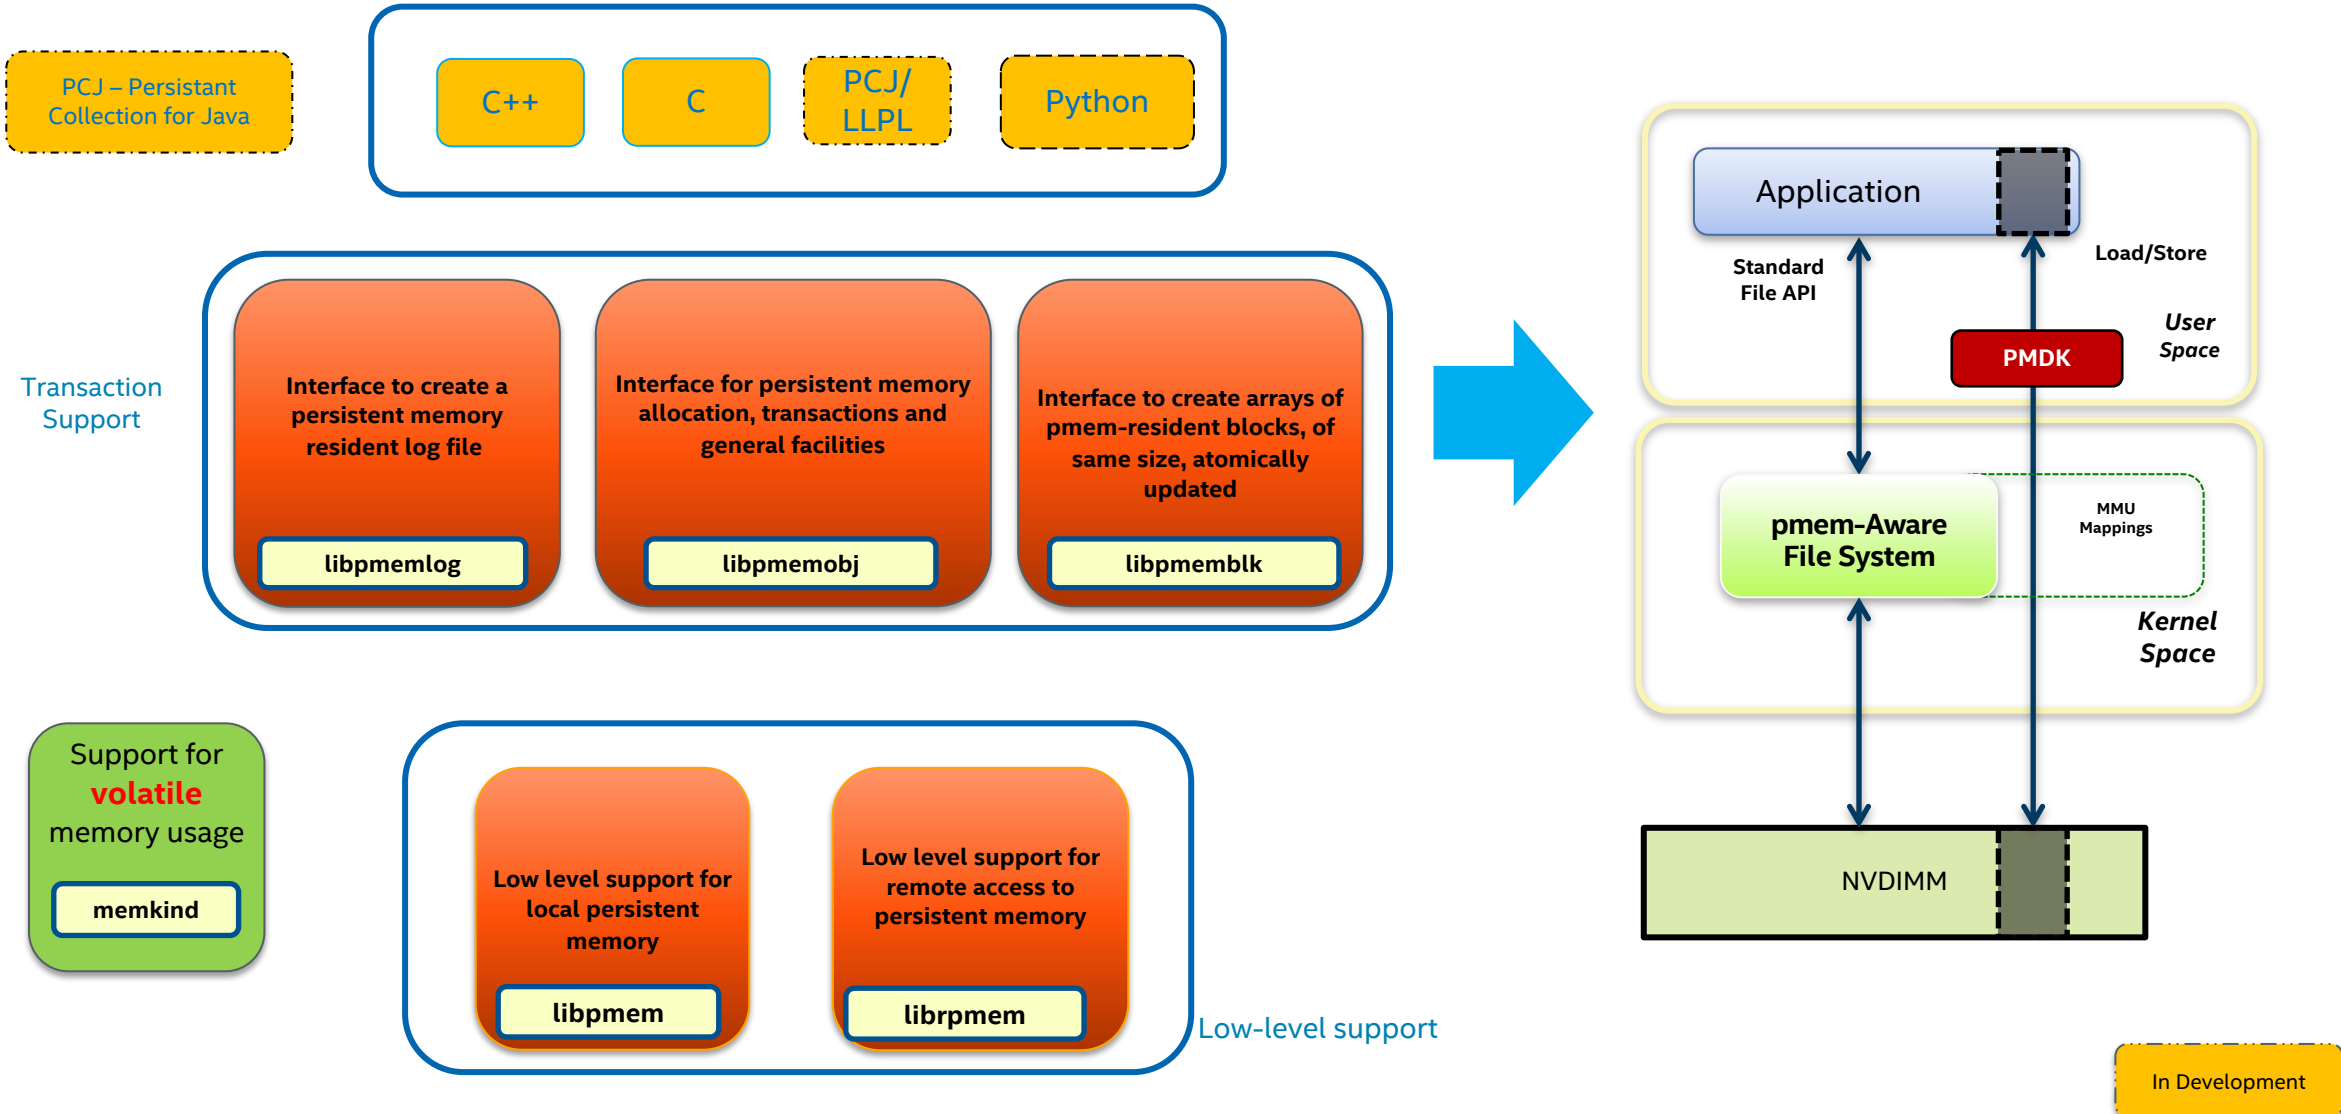
The Persistent Memory Development Kit

PMDK <http://pmem.io>



- PMDK is a collection of libraries
 - Developers pull only what they need
 - Low level programming support
 - Transaction APIs
 - Fully validated
 - Performance tuned.
- Open Source & Product neutral

PMDK Libraries



ESSENTIAL LIBP MEM KNOWLEDGE

libpmem examples

Source: <https://github.com/pmem/pmdk/tree/master/src/examples/libpmem>

```
/*
 * simple_copy.c -- show how to use pmem_memcpy_persist()
 *
 * usage: simple_copy src-file dst-file
 *
 * Reads 4k from src-file and writes it to dst-file.
 */

/* create a pmem file and memory map it */
if ((pmemaddr = pmem_map_file(argv[2], BUF_LEN,
                             PMEM_FILE_CREATE|PMEM_FILE_EXCL,
                             0666, &mapped_len, &is_pmem)) == NULL) {
    perror("pmem_map_file");
    exit(1);
}
```

Using is_pmem

```
if (is_pmem) {  
    pmem_memcpy_persist(pmemaddr, buf, cc);  
} else {  
    memcpy(pmemaddr, buf, cc);  
    pmem_msync(pmemaddr, cc);  
}
```


ESSENTIAL LIBPMEMBLK KNOWLEDGE

libpmemblk examples

Source: <https://github.com/pmem/pmdk/blob/master/src/examples/libpmemblk>

```
/*
 * manpage.c – simple example for libpmemblk manpage
 */

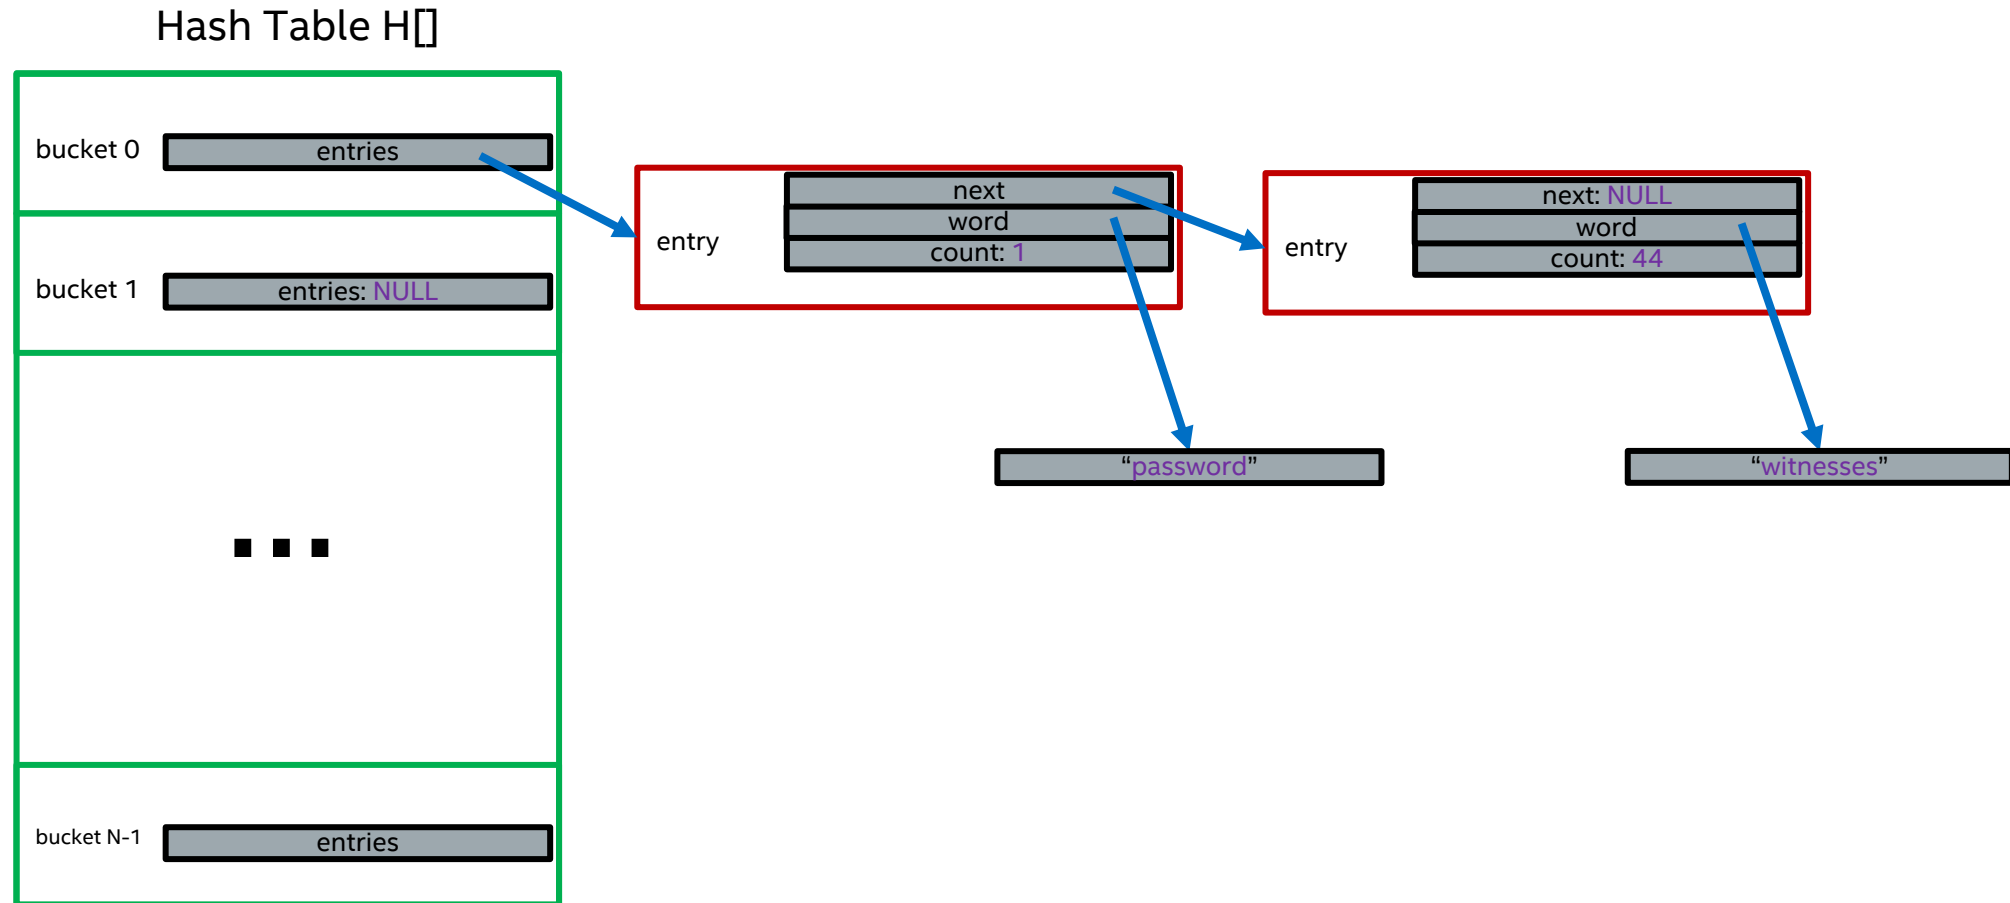
/* store a block at index 5 */
strcpy(buf, "hello, world");
if (pmemblk_write(pbp, buf, 5) < 0) {
    perror("pmemblk_write");
    exit(1);
}

/* read the block at index 10 (reads as zeros initially) */
if (pmemblk_read(pbp, buf, 10) < 0) {
    perror("pmemblk_read");
    exit(1);
}

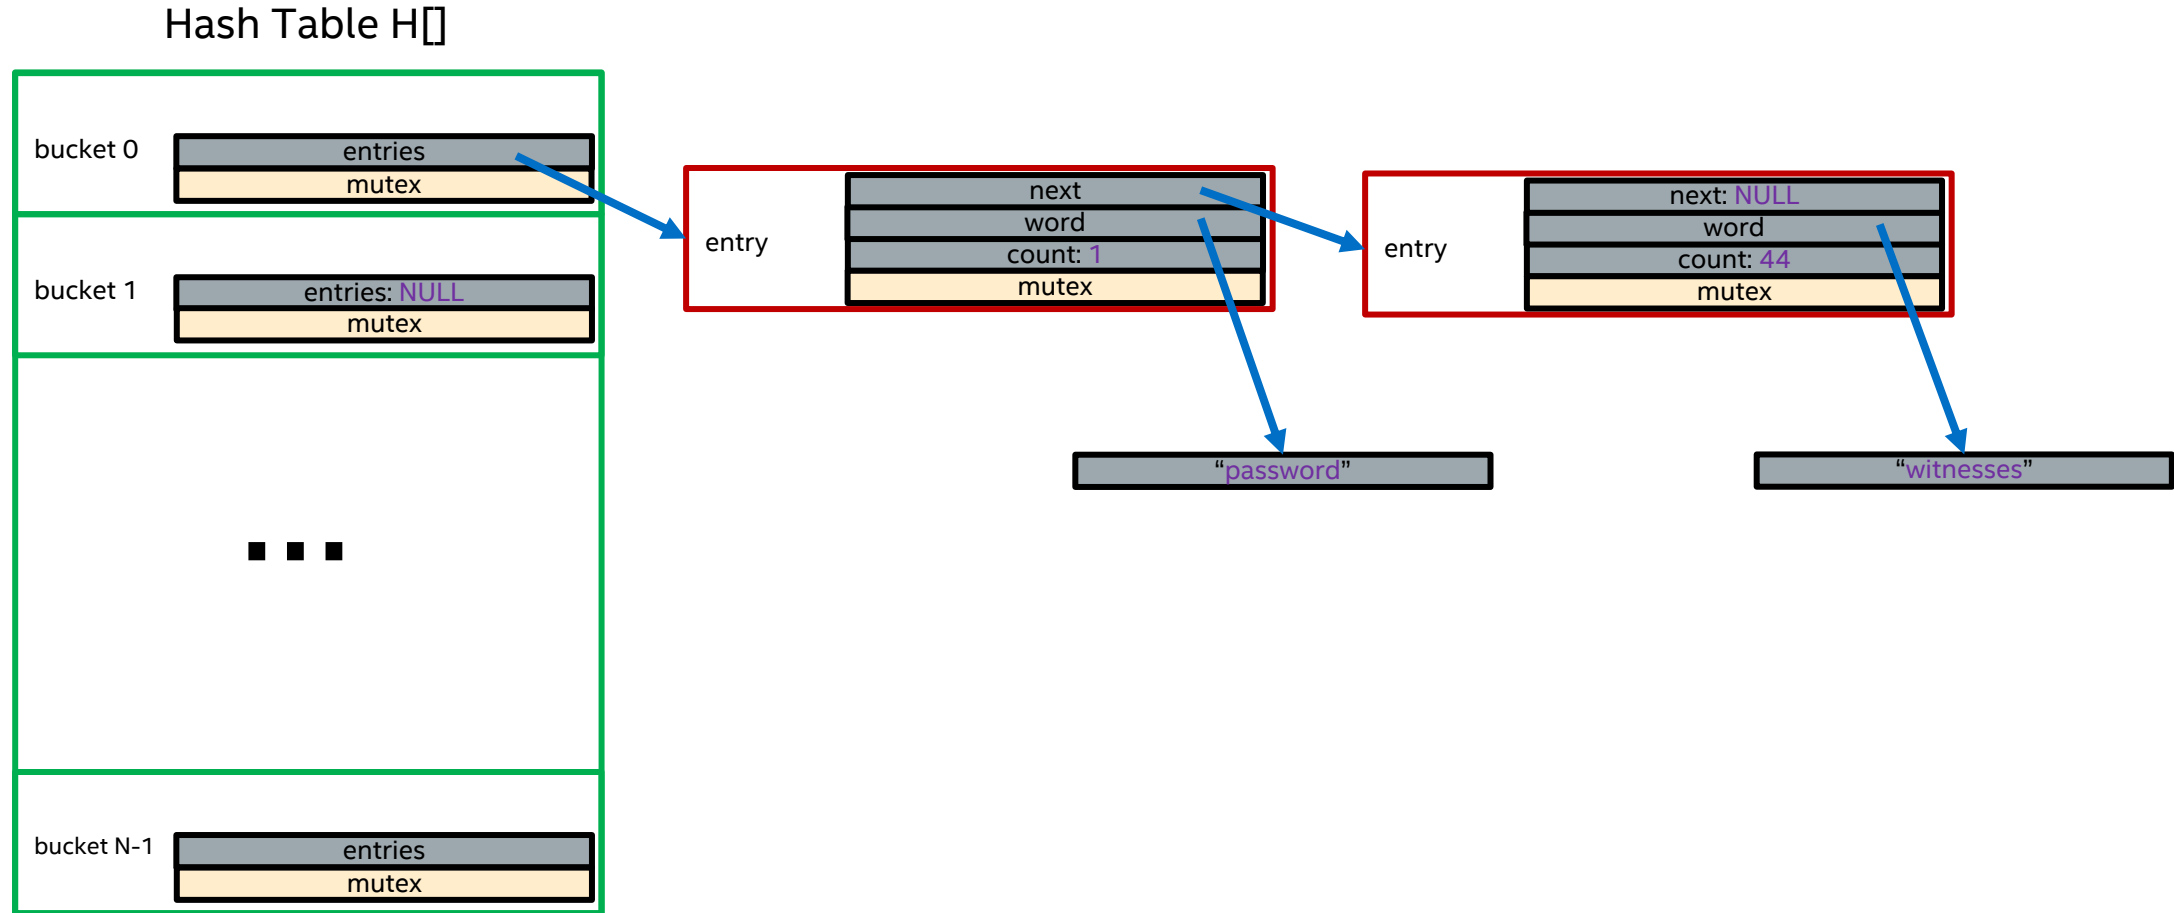
/* zero out the block at index 5 */
if (pmemblk_set_zero(pbp, 5) < 0) {
    perror("pmemblk_set_zero");
    exit(1);
}
```

LIBPMEMOBJ EXAMPLE


Simple C program to build example on (nothing related to pmem yet)



Adding multi-threading support (nothing related to pmem yet)




The *Root Object*



root object:


- assume it is always there
- created first time accessed
- initially zeroed

Using the Root Object

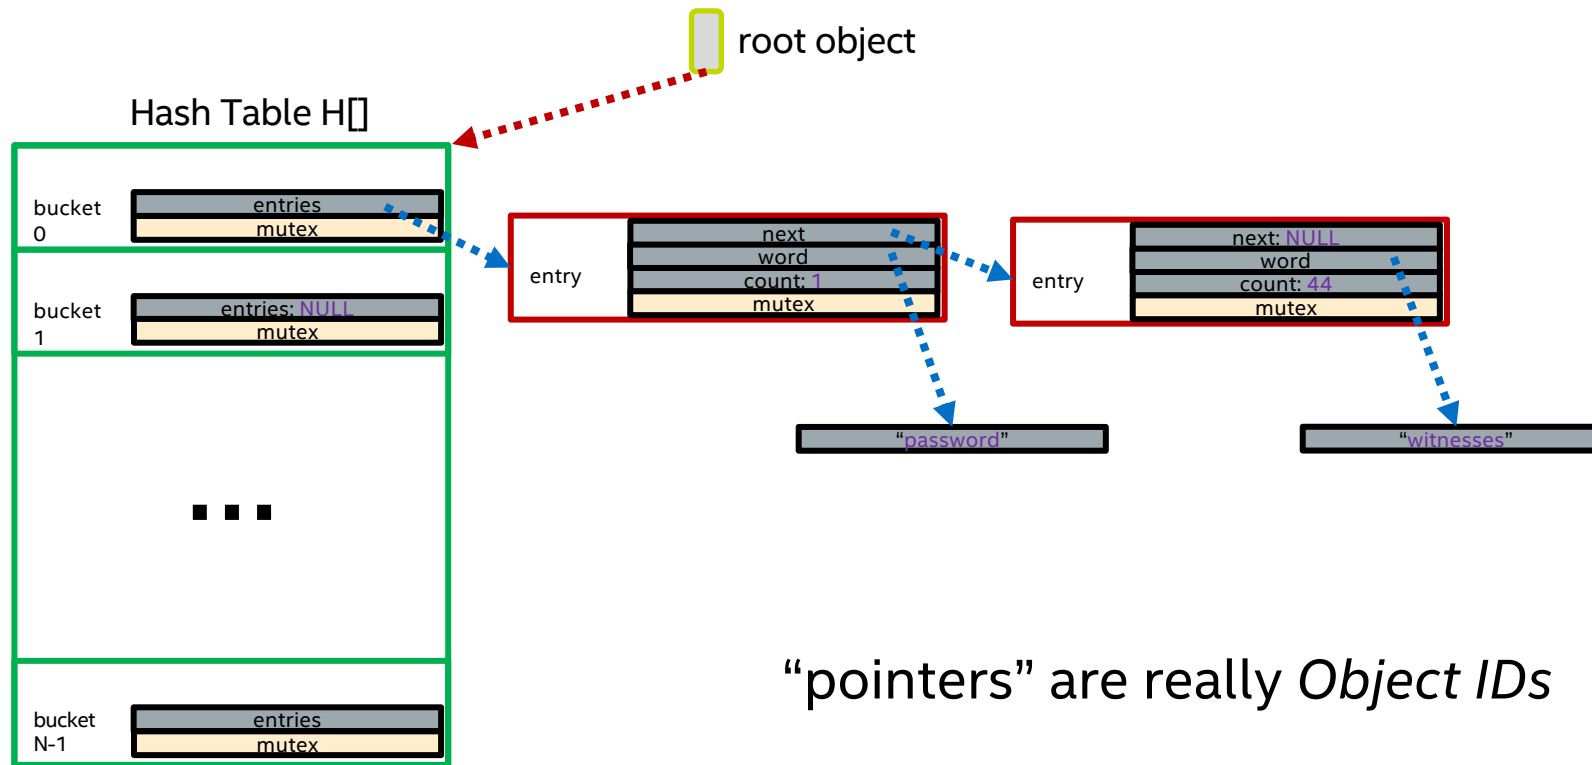


Link pmem data structures in pool off the root object to find them on each program run

“pointers” are really *Object IDs*



Moving data the example to pmem



C Programming with libpmemobj

Transaction Syntax

```
TX_BEGIN(Pop) {  
    /* the actual transaction code goes here... */  
} TX_ONCOMMIT {  
    /*  
     * optional - executed only if the above block  
     * successfully completes  
     */  
} TX_ONABORT {  
    /*  
     * optional - executed if starting the transaction fails  
     * or if transaction is aborted by an error or a call to  
     * pmemobj_tx_abort()  
     */  
} TX_FINALLY {  
    /*  
     * optional - if exists, it is executed after  
     * TX_ONCOMMIT or TX_ONABORT block  
     */  
} TX_END /* mandatory */
```

Properties of Transactions

Powerfail
Atomicity

Multi-
Thread
Atomicity

```
TX_BEGIN_PARAM(Pop, TX_PARAM_MUTEX, &D_RW(ep)->mtx, TX_PARAM_NONE) {  
    TX_ADD(ep);  
    D_RW(ep)->count++;  
} TX_END
```

Caller must
instrument code
for undo logging

Persistent Memory Locks

- Want locks to live near the data they protect (i.e. inside structs)
- Does the state of locks get stored persistently?
 - Would have to flush to persistence when used
 - Would have to recover locked locks on start-up
 - Might be a different program accessing the file
 - Would run at pmem speeds
- PMEMmutex
 - Runs at DRAM speeds
 - Automatically initialized on pool open

C++ Programming with libpmemobj

C++ Queue Example: Declarations

```
/* entry in the queue */
struct pmem_entry {
    persistent_ptr<pmem_entry> next;
    p<uint64_t> value;
};
```

<code>persistent_ptr<T></code>	Pointer is really a position-independent Object ID in pmem. Gets rid of need to use C macros like <code>D_RW()</code>
<code>p<T></code>	Field is pmem-resident and needs to be maintained persistently. Gets rid of need to use C macros like <code>TX_ADD()</code>

C++ Queue Example: Transaction

```
void push(pool_base &pop, uint64_t value) {  
    transaction::run(pop, [&] {  
        auto n = make_persistent<pmem_entry>();  
  
        n->value = value;  
        n->next = nullptr;  
        if (head == nullptr) {  
            head = tail = n;  
        } else {  
            tail->next = n;  
            tail = n;  
        }  
    });  
};
```

Transactional
(including allocations &
frees)

PCJ EXAMPLE

Persistent Containers for Java

Library of persistent classes

- object state stored on a persistent heap
- stored in object layout form, no serialization or deserialization
- instances behave like regular Java objects, just longer-lived
- reachability-based lifetime
- easy-to-understand data consistency model

API for defining persistent classes

- expressiveness similar to that of regular classes
- no change to developer toolchain

Separate library for low-level access to pmem

- byte-addressable persistent memory regions
- developer can roll their own abstractions

<https://github.com/pmem/pcj>

Persistent Classes

Primitive arrays (e.g. PersistentByteArray, mutable and immutable)
PersistentArray<E extends AnyPersistent> (mutable and immutable)
PersistentTuple<T1 extends AnyPersistent, ...> (mutable and immutable)
PersistentArrayList<E extends AnyPersistent>
PersistentHashMap<K extends AnyPersistent, V extends AnyPersistent>
PersistentLinkedList<E extends AnyPersistent>
PersistentLinkedQueue<E extends AnyPersistent>
PersistentSkipListMap<K extends AnyPersistent, V extends AnyPersistent>
PersistentFPTreeMap<K extends AnyPersistent, V extends AnyPersistent>
PersistentSIHashMap<K extends AnyPersistent, V extends AnyPersistent>
ObjectDirectory - indefinitely reachable root map of <String, T extends AnyPersistent>

Primitive types (as field and array element values, no separate class)
Boxed primitives (e.g. PersistentLong)
PersistentString
PersistentByteBuffer
PersistentUUID
PersistentAtomicReference<T extends AnyPersistent>

WordFrequency.java

```
public class WordFrequency {  
    private static Map<String, Integer> counts = new TreeMap<>();  
  
    public static void main(String[] args) {  
        if (args.length == 0) System.out.println("usage: WordFrequency <list of files to process>");  
        for (int i = 0; i < args.length; i++) {  
            try {  
                Scanner scanner = new Scanner(new File(args[i]));  
                while (scanner.hasNext()) {  
                    String word = scanner.next();  
                    counts.merge(word, 1, Integer::sum);  
                }  
            }  
            catch (FileNotFoundException fnf) {throw new RuntimeException(fnf.getCause());}  
        }  
  
        // print counts  
        for (Map.Entry<String, Integer> e : counts.entrySet()) {  
            System.out.format("%d %s\n", e.getValue().intValue(), e.getKey());  
        }  
    }  
}
```

ParallelWordFrequency.java

```
public class ParallelWordFrequency {
    private static Map<String, Integer> counts = new ConcurrentSkipListMap<>();

    public static void main(String[] args) throws InterruptedException {
        if (args.length == 0) System.out.println("usage: ParallelWordFrequency <list of files to process>");
        Thread[] ts = new Thread[args.length];
        for (int i = 0; i < args.length; i++) {
            int ii = i;
            ts[ii] = new Thread(() -> {
                try {
                    Scanner scanner = new Scanner(new File(args[ii]));
                    while (scanner.hasNext()) {
                        String word = scanner.next();
                        counts.merge(word, 1, Integer::sum);
                    }
                }
                catch (FileNotFoundException fnf) {throw new RuntimeException(fnf.getCause());}
            });
        }
        for (Thread t : ts) t.start();
        for (Thread t : ts) t.join();
    }
}
```

...

PersistentParallelWordFrequency.java

```
package examples.wordfrequency;

import java.io.File;
import java.io.FileNotFoundException;
import java.util.Scanner;
import lib.util.persistent.PersistentHashMap;
import lib.util.persistent.PersistentSkipListMap;
import lib.util.persistent.PersistentString;
import lib.util.persistent.PersistentInteger;
import lib.util.persistent.ObjectDirectory;
import java.util.Map;
import java.util.function.BiFunction;
```

PersistentParallelWordFrequency.java

```
public class PersistentParallelWordFrequency {  
    private static Map<PersistentString, PersistentInteger> counts = getOrInitializeCounts();  
  
    @SuppressWarnings("unchecked")  
    private static Map<PersistentString, PersistentInteger> getOrInitializeCounts() {  
        String DATA_KEY = "WordFrequencyData";  
        PersistentSkipListMap<PersistentString, PersistentInteger> map =  
            ObjectDirectory.get(DATA_KEY, PersistentSkipListMap.class);  
        if (map == null) ObjectDirectory.put(DATA_KEY, map = new PersistentSkipListMap<>());  
        return map;  
    }  
}
```

PersistentParallelWordFrequency.java

```
public static void main(String[] args) throws InterruptedException {
    if (args.length == 0) System.out.println("usage: PersistentParallelWordFrequency <list of files to process>");
    final PersistentInteger ONE = new PersistentInteger(1);
    Thread[] ts = new Thread[args.length];
    for (int i = 0; i < args.length; i++) {
        int ii = i;
        ts[ii] = new Thread(() -> {
            try {
                Scanner scanner = new Scanner(new File(args[ii]));
                while (scanner.hasNext()) {
                    PersistentString word = new PersistentString(scanner.next());
                    counts.merge(word, ONE, PersistentParallelWordFrequency::sum);
                }
            }
            catch (FileNotFoundException fnf) {throw new RuntimeException(fnf.getCause());}
        });
    }
    for (Thread t : ts) t.start();
    for (Thread t : ts) t.join();
}
```

PersistentParallelWordFrequency.java

```
// print current counts
    for (Map.Entry<PersistentString, PersistentInteger> e : counts.entrySet()) {
        System.out.format("%d %s\n", e.getValue().intValue(), e.getKey());
    }
}

public static PersistentInteger sum(PersistentInteger x, PersistentInteger y) {
    return new PersistentInteger(x.intValue() + y.intValue());
}
}
```


LINKS TO MORE INFORMATION

Using Persistent Memory as Volatile Memory

Bigger/cheaper than DRAM

Application decides what lives in DRAM, what lives in persistent memory

- Unlike Memory Mode, where HW decides

Similar to NUMA programming

- app allocates different “kinds” of memory

memkind library: <http://memkind.github.io/memkind/>

- Familiar malloc/free style programming with multiple pools
 - NUMA nodes, HBM, etc.
- Can construct pools with persistent memory

libpmemkv

<https://github.com/pmem/pmemkv>

General-purpose key-value store

- Simple API, handles pmem transactions, etc so caller doesn't need to
- Multiple storage engines, tuned for pmem
- Multiple language bindings: C, C++, Java, Ruby, JavaScript

Still “experimental” – in the process of validating to product quality

More Developer Resources

- Find the PMDK (Persistent Memory Development Kit) at <http://pmem.io/pmdk/>
- Getting Started
 - Intel IDZ persistent memory- <https://software.intel.com/en-us/persistent-memory>
 - Entry into overall architecture - <http://pmem.io/2014/08/27/crawl-walk-run.html>
 - Emulate persistent memory - <http://pmem.io/2016/02/22/pm-emulation.html>
- Linux Resources
 - Linux Community Pmem Wiki - <https://nvdimm.wiki.kernel.org/>
 - Pmem enabling in SUSE Linux Enterprise 12 SP2 - <https://www.suse.com/communities/blog/nvdimm-enabling-suse-linux-enterprise-12-service-pack-2/>
- Windows Resources
 - Using Byte-Addressable Storage in Windows Server 2016 - <https://channel9.msdn.com/Events/Build/2016/P470>
 - Accelerating SQL Server 2016 using Pmem - <https://channel9.msdn.com/Shows/Data-Exposed/SQL-Server-2016-and-Windows-Server-2016-SCM--FAST>
- Other Resources
 - SNIA Persistent Memory Summit 2018 - <https://www.snia.org/pm-summit>
 - Intel manageability tools for Pmem - <https://01.org/ixpdimm-sw/>

Intel Developer Support & TOOLS

PMDK Tools

- Valgrind plugin: pmemcheck
- Debug mode, tracing, pmembench, pmreorder

pmem.io

New features to support Intel® Optane™ DC persistent memory

- Intel® VTune™ Amplifier – Performance Analysis
- Intel® Inspector – Persistence Inspector finds missing cache flushes & more
- Free downloads available

software.intel.com/pmem

Q&A